

TYPEWHICH Guide

Luna Phipps-Costin, Carolyn Jane Anderson, Michael Greenberg, and Arjun Guha

April 20, 2022

Contents

1	Introduction	5
1.1	Building and Testing TYPEWHICH	5
2	Artifact Evaluation: Getting Started	9
3	Artifact Evaluation: Step by Step Guide	11
3.1	Claims To Validate	11
3.1.1	GTLC Benchmarks on Multiple Tools (Figure 15)	11
3.1.2	Grift Benchmarks with TYPEWHICH	12
3.1.3	Performance	13
3.2	Exploring Type Migrations	13
3.3	Input Language	14
4	Guide to Source Code	15

Chapter 1

Introduction

TYPEWHICH is a type migration tool for the gradually-typed lambda calculus with several extensions. Its distinguishing characteristics are the following:

1. TYPEWHICH formulates type migration as a MaxSMT problem.
2. TYPEWHICH always produces a migration, as long as the input program is well-scoped.
3. TYPEWHICH can optimize for different properties: it can produce the most informative types, or types that ensure compatibility with un-migrated code.

Before you read the rest of this guide or try to use TYPEWHICH, we strongly recommend reading Phipps-Costin et al. (2021), which describes TYPEWHICH in depth.

This repository contains the source code for TYPEWHICH. In addition to the core type migration algorithm, the TYPEWHICH executable has several auxiliary features:

1. It has a parser for the Grift programming language, which we use to infer types for the Grift benchmarks from Kuhlenschmidt et al. (2019);
2. It has an interpreter for the GTLC, which we use in validation;
3. It has an implementation of the gradual type inference algorithm from Rastogi et al. (2012); and
4. It includes a framework for evaluating type migration algorithms, which we use to compare TYPEWHICH to several algorithms from the literature Rastogi et al. (2012); Campora et al. (2018); Migeed and Palsberg (2020); Siek and Vachharajani (2008).

Finally, this repository contains several gradual typing benchmarks:

1. The “challenge set” from Phipps-Costin et al. (2021);
2. The benchmarks from Migeed and Palsberg (2020); and
3. The benchmarks from Kuhlenschmidt et al. (2019).

This document will guide you through building TYPEWHICH, using it on example programs, and using the evaluation framework to reproduce our experimental results.

1.1 Building and Testing TYPEWHICH

For artifact evaluation, we strongly recommend using the TYPEWHICH Virtual Machine and skipping this section.

TYPEWHICH is built in Rust and uses Z3 under the hood. In principle, it should work on macOS, Linux or Windows, though we have only tried it on macOS and Linux. *However*, our evaluation uses the implementation from Siek and Vachharajani (2008), which is an old piece of software that is difficult to build on a modern platform. We have managed to compile it a Docker container and produce a 32-bit Linux binary. It should be possible to build it for other platforms, but it will require additional effort. Therefore, **we strongly recommend using Linux to evaluate TYPEWHICH.**

Installing TYPEWHICH Dependencies To build TYPEWHICH from source, you will need:

1. The Rust language toolchain.
2. The Z3 build dependencies and the “usual” build toolchain. On Ubuntu Linux, you can run the following command to get them:

```
sudo apt-get install libz3-dev build-essential
```

3. Python 3 and PyYAML to run the integration tests. These are installed by default on most platforms. If you can run the following command successfully then you already have them installed:

```
python3 -c "import yaml"
```

Installing Other Type Migration Tools TYPEWHICH does not require these dependencies, but they are necessary to reproduce our evaluation.

1. Migeed and Palsberg (2020) is implemented in Haskell. We have written a parser and printer for their tool that is compatible with TYPEWHICH. This modified implementation is available at the following URL:

<https://github.com/arjunguha/migeed-palsberg-popl2020>

Build the tool as described in the repository, and then copy (or symlink) the `MaxMigrate` program to `bin/MaxMigrate` in the TYPEWHICH directory. On Linux, the executable is at:

```
migeed-palsberg-popl2020/.stack-work/install/x86_64-linux-tinfo6/  
lts-13.25/8.6.5/bin/MaxMigrate
```

2. Siek and Vachharajani (2008) is implemented in OCaml 3.12 (which is quite old). The following repository has an implementation of the tool, with a modified parser and printer that is compatible with TYPEWHICH:

<https://github.com/arjunguha/siek-vachharajani-dls2008>

Build the tool as described in the repository, and then copy (or symlink) the `gtlc` program to `bin/gtubi` in the TYPEWHICH directory.

Warning: The repository builds a 32-bit Linux executable. You will need to ensure that your Linux system has the libraries needed to run 32-bit code.

3. Campora et al. (2018) The following repository has our implementation of the algorithm from Campora et al. (2018):

<https://github.com/arjunguha/mgt>

Build the tool as described in the repository and then copy (or symlink) the `mgt` program to `bin/mgt` in the TYPEWHICH directory.

Note: The original implementation by the authors of Campora et al. (2018) does not produce an ordinary migrated program as output. Instead, it produces a BDD that can be interpreted as a family of programs. Our implementation of their algorithm produces programs as output.

Building and Testing Use cargo to build TYPEWHICH:

```
cargo build
```

Run the unit tests:

```
cargo test
```

You may see a few ignored tests, but *no tests should fail*.

Test TYPEWHICH using the Grift benchmarks:

```
./test-runner.sh grift grift
```

No tests should fail.

Finally, run the GTLC benchmarks without any third-party tools:

```
cargo run -- benchmark benchmarks.yaml \  
--ignore Gtubi MGT MaxMigrate > test.results.yaml
```

You will see debugging output (on standard error), but the results will be saved to the YAML file. Compare these results to known good results:

```
./bin/yamldiff test.expected.yaml test.results.yaml
```

You should see no output, which indicates that there are no differences.

Build TYPEWHICH in release mode (only needed for performance evaluation):

```
cargo build --release
```


Chapter 2

Artifact Evaluation: Getting Started

Before starting this chapter, we must either:

- Use the TYPEWHICH Virtual Machine, or
- Install TYPEWHICH manually, along with all the third party tools we use for evaluation.

Warning: TYPEWHICH uses the Z3 SMT solver under the hood, and different versions of Z3 can produce slightly different results. The expected outputs that we document in this guide were produced on the TYPEWHICH Virtual Machine.

If the following steps are successful, then we can be quite confident that TYPEWHICH and all third-party tools are working as expected.

1. From a terminal window, enter the TYPEWHICH directory:

```
cd ~/typewhich
```

2. Run the TYPEWHICH benchmarks and output results to `results.yaml`:

```
./bin/TypeWhich benchmark benchmarks.yaml > results.yaml
```

This will take less than five minutes to complete. This command runs the GTLC benchmarks using all five tools, including TYPEWHICH in two modes. Therefore, for each benchmark, we will see six lines of output (on standard error). For example:

```
Running Gtubi on adversarial/01-farg-mismatch.gtlc ...
Running InsAndOuts on adversarial/01-farg-mismatch.gtlc ...
Running MGT on adversarial/01-farg-mismatch.gtlc ...
Running MaxMigrate on adversarial/01-farg-mismatch.gtlc ...
Running TypeWhich2 on adversarial/01-farg-mismatch.gtlc ...
Running TypeWhich on adversarial/01-farg-mismatch.gtlc ...
```

There are three runs of third-party tools that take longer than 30 seconds, so you will `Killed` appear three times. These are known shortcomings that are described in the paper.

3. Run the following command to ensure that the results are identical to known good results:

```
./bin/yamldiff expected.yaml results.yaml
```

There should be no output printed, which indicates that there are no differences.

4. Run the following command to run TYPEWHICH on the Grift benchmarks:

```
./grift_inference.sh
```

We expect to see `MATCHES` print several times, which indicates that TYPEWHICH inferred exactly the same types that were written by the Grift authors on that benchmark. However, we also expect to see a `Warning`, and two mismatches on `n_body` and `sieve`.

At this point, we can investigate the artifact in more depth, which is the subject of the next chapter.

Chapter 3

Artifact Evaluation: Step by Step Guide

This chapter assumes you have completed the steps in Chapter 2.

3.1 Claims To Validate

The paper makes the following claims that we validate in this chapter:

1. Figure 15 summarizes the performance of several type migration tools on a suite of benchmarks. This artifact generates that figure, and we can validate the data and benchmarking scripts in as much depth as desired.
2. Section 6.5 runs TYPEWHICH on benchmarks written in Grift. These benchmarks have two versions: one that has no type annotations, and the other that has human-written type annotations. When run on the unannotated Grift benchmarks, TYPEWHICH recovers the human-written annotations on all but two of the Grift benchmarks. This artifact includes a script that produces this result.
3. Section 6.6 reports that our full suite of benchmarks is 892 LOC, and TYPEWHICH takes three seconds to run on all of them. This artifact includes the script we use for the performance evaluation. It will take longer in a virtual machine, but should be roughly the same. i.e., it will be significantly less than 30 seconds.

The rest of this section will be a step-by-step guide through repeating and validating these claims.

3.1.1 GTLC Benchmarks on Multiple Tools (Figure 15)

In the previous chapter, we generated `results.yaml`. That ran TYPEWHICH and all other tools on two suites of benchmarks:

1. All the benchmarks from Migeed and Palsberg (2020), which are in the `migeed` directory.
2. The “challenge set” from the paper, which are in the `adversarial` directory.

The file `benchmarks.yaml` drives the benchmarking framework. The top of the file lists the type migration tool, and is followed by a list of benchmark files, and some additional information that needed to produce results. The entire benchmarking procedure is implemented in `src/benchmark.rs`, which does performs the following steps on each benchmark:

1. It checks that the tool produces valid program, to verify that the tool did not reject the program.
2. It runs the original program and the output of the tool and checks that they produce the same result, to verify that the tool did not introduce a runtime error.

3. In a gradually typed language, increasing type precision can make a program incompatible with certain contexts. To check if this is the case, every benchmark in the YAML file *may* be accompanied by a context that witnesses the incompatibility: the framework runs the original and migrated program in the context, to check if they produce different results.
4. The framework counts the number of anys that are eliminated by the migration tool. Every eliminated any improves precision, but *may or may not* introduce an incompatibility, but this requires human judgement. For example, in the program `fun x . x + 1`, annotating “x” with `int` does not introduce an incompatibility. However, in `fun x . x`, annotating “x” with `int` is an incompatibility. The framework flags these results for manual verification. However, it allows the input YAML to specify expected outputs to suppress these warnings when desired.

The file `results.yaml` is a copy of `benchmarks.yaml` with output data added by the benchmarking framework. We use this file to generate Figure 15 in the paper. You should validate that table as follows:

1. Check that `results.yaml` does not have any errors: look for the string “Disaster” in that file. It should not occur!
2. Regenerate the LaTeX snippet for the table with the following command:

```
./bin/TypeWhich latex-benchmark-summary results.yaml
```

The output that you will see is roughly the LaTeX code for Figure 15, with two small differences:

- (a) It prints `TypeWhich2` instead of `TypeWhichC`, and
- (b) `TypeWhich` instead of `TypeWhichP`.

However, the order of rows and columns is exactly the same as the table in the paper. It should be straightforward to check that the fractions in this output are exactly the fractions reported in the table.

3.1.2 Grift Benchmarks with TYPEWHICH

The Grift evaluation script (`grift_inference.sh`) uses the `-compare` flag of `TYPEWHICH`, which corresponds migrated types to the provided file’s type annotations and reports whether they match, ignoring annotations, coercions, and unannotated identifiers.

On benchmarks for which it reports `MATCHES`, `TYPEWHICH` produced exactly the same type annotations as the hand-typed versions.

On `n_body`, verify that `grift-suite/benchmarks/src/dyn/n_body.grift` and `grift-suite/benchmarks/src/dyn/n_body_no_unused_funs.grift` differ only by the removal of unused getters and setters near the top of the program. Note that `TYPEWHICH`’s types on the adjusted benchmark with no unused functions matches the hand-typed version.

On `sieve`, the warning refers to a lack of parsing support for recursive types. As a result the mismatch message is less informative than inspection. To verify exactly which types fail to migrate, run `TYPEWHICH` to migrate the types of the program:

```
./bin/TypeWhich migrate grift-suite/benchmarks/src/dyn/sieve.grift
```

Each bound identifier (excluding lets) will be printed, with its type. Keeping that input open, manually inspect the hand-typed version at `grift-suite/benchmarks/src/static/sieve/single/sieve.grift`. Consider, for example, the first identifier, `stream-first`. The annotated program declares `stream-first` to accept `(Rec s (Tuple Int (-> s)))` and return `Int`, while `TYPEWHICH`’s output accepts any and returns any. Inspecting each function remaining, you will see that every `(Rec s (Tuple Int (-> s)))` is replaced with the dynamic type. Also, some (but not all) integer types are migrated as the dynamic type (because they hold values from projections out of tuples of any). Note that the unit-terminated pair representation of tuples is visible in `stream-unfold`, which otherwise has the expected type.

3.1.3 Performance

From the TYPEWHICH directory, run the following command:

```
time ./performance.sh
```

The script will take roughly three seconds to complete. You can read the script to verify that it runs TYPEWHICH on three suites of benchmarks:

1. `migeed/*.gtlc`: the benchmarks from Migeed and Palsberg (2020),
2. `adversarial/*.gtlc`: the “challenge set” from our paper, and
3. `grift-suite/benchmarks/src/dyn/*.grift`: the benchmarks from Kuhlenschmidt et al. (2019).

3.2 Exploring Type Migrations

Our artifact includes several type migration tools, in addition to TYPEWHICH, and we have hacked their parsers to work with the same concrete syntax, so that it is easy to use any tool on the same program. We encourage you to try some out, and to modify the benchmarks as well. Here are the available tools:

- To run Migeed and Palsberg (2020):

```
./bin/MaxMigrate FILENAME.gtlc
```

- To run Campora et al. (2018):

```
./bin/mgt FILENAME.gtlc
```

- To run Siek and Vachharajani (2008):

```
./bin/gtubi FILENAME.gtlc
```

- To run Rastogi et al. (2012):

```
./bin/TypeWhich migrate --ins-and-outs FILENAME.gtlc
```

- To run TYPEWHICH and produce types that are safe in all contexts:

```
./bin/TypeWhich migrate FILENAME.gtlc
```

- To run TYPEWHICH and produce precise types that may not work in all contexts:

```
./bin/TypeWhich migrate --precise FILENAME.gtlc
```

Example Create a file called `input.gtlc` with the following contents:

```
(fun f. (fun y. f) (f 5)) (fun x. 10 + x)
```

This program omits all type annotations: TYPEWHICH assumes that omitted annotations are all **any**. We can migrate the the program using TYPEWHICH in two modes:

1. In *compatibility mode*, TYPEWHICH infers types but maintains compatibility with un-migrated code:

```
$ ./bin/TypeWhich migrate input.gtlc
(fun f:any -> int. (fun y:int. f) (f 5)) (fun x:any. 10 + x)
```

2. In *precise mode*, TYPEWHICH infers the most precise type that it can, though that may come at the expense of compatibility:

```
$ ./bin/TypeWhich migrate --precise input.gtlc
(fun f:int -> int. (fun y:int. f) (f 5)) (fun x:int. 10 + x)
```

3.3 Input Language

TYPEWHICH supports a superset of the GTLC, written in the following syntax. Note that the other tools do not support all the extensions documented below.

b	<code>:= true false</code>	Boolean literal
n	<code>:= ... -1 0 1 ...</code>	Integer literals
s	<code>:= "..."</code>	String literals
c	<code>:= b n s</code>	Literals
T	<code>:= any</code>	The unknown type
	<code> int</code>	Integer type
	<code> bool</code>	Boolean type
	<code> $T_1 \rightarrow T_2$</code>	Function type
	<code> (T)</code>	
e	<code>:= x</code>	Bound identifier
	<code> c</code>	Literal
	<code> e : T</code>	Type ascription
	<code> (e)</code>	Parenthesis
	<code> fun x . e</code>	Function
	<code> $e_1 e_2$</code>	Application
	<code> $e_1 + e_2$</code>	Addition
	<code> $e_1 * e_2$</code>	Multiplication
	<code> $e_1 = e_2$</code>	Integer equality
	<code> $e_1 +? e_2$</code>	Addition or string concatenation (overloaded)
	<code> (e_1, e_2)</code>	Pair
	<code> fix f . e</code>	Fixpoint
	<code> if e_1 then e_2 else e_3</code>	Conditional
	<code> let $x = e_1$ in e_2</code>	Let binding
	<code> let rec $x = e_1$ in e_2</code>	Recursive let binding

Chapter 4

Guide to Source Code

The root TypeWhich directory includes a number of utilities, programs, and source code (though most of TYPEWHICH is provided in `src/`):

1. `adversarial/`, `grift-suite/benchmarks/`, `migeed/`: The three components of the TYPEWHICH benchmark suite, `adversarial/` being original, and `grift-suite/` and `migeed/` adapted from the referenced research
2. `doc/`: Source and render of this documentation
3. `benchmarks.yaml`: This is the test harness configuration and data for the TYPEWHICH benchmarks framework, specifying to run the benchmarks and tools presented in the paper
4. `expected.yaml`: Provides the expected behavior of the tool when configured with `benchmarks.yaml`
5. `test.expected.yaml`: Provides the expected behavior of only TYPEWHICH/ Rastogi et al. (2012) for testing the implementations
6. `bin/`: Provides (and expects user to provide) symbolic links to tools
7. `build.rs`, `Cargo.lock`, `Cargo.toml`, `target/`: Required build files for TYPEWHICH. Binaries are placed in `target/`
8. `other-examples/`: Provides additional programs that are not interesting enough to be in the TYPEWHICH benchmark suite
9. `grift_inference.sh`: Evaluation tool for Grift benchmarks which compares if types produced are exactly the same as the static types provided in the suite
10. `performance.sh`, `test-runner.sh`, `run_tool.sh`: Tools that run TYPEWHICH on more programs or in release mode
11. **`src/`**: The TYPEWHICH implementation, including implementation of Rastogi et al. (2012)

Within `src/`, the following files are found:

1. **`benchmark.rs`**, **`precision.rs`**: Provides the TYPEWHICH benchmarking framework
2. **`cgen.rs`**: Generates the documented constraints of the TYPEWHICH algorithm and performs type migration
3. **`eval.rs`**: An interpreter for the GTLC with explicit coercions

4. `insert_coercions.rs`: Type-directed coercion insertion for the GTLC, used for the interpreter. Not related to type migration
5. `grift.l`, `grift.y`, `grift.rs`, `lexer.l`, `parser.y`, `pretty.rs`: Parsers and printers for Grift and the unified concrete syntax used by all tools
6. `ins_and_outs/`: Our implementation of Rastogi et al. (2012).
7. `main.rs`: Entry point; options parsing
8. `syntax.rs`: The language supported by TYPEWHICH. Also includes the `-compare` tool used for Grift evaluation
9. `type_check.rs`: Type-checking for programs with explicit coercions
10. `z3_state.rs`: Abstraction for the Z3 solver used for type inference in TYPEWHICH

The core of the TYPEWHICH algorithm is found in `cgen.rs`. The constraints specified in the paper are implemented in `State::cgen` (~line 52), with comments resembling the notation from the paper. Of note are references to `strengthen` and `weaken`, which are simply macros for $(t1 = t2 \wedge w) \vee (t1 = * \wedge \text{ground}(t2) \wedge \neg w)$, w fresh; and $(t1 = t2 \wedge w) \vee (t2 = * \wedge \text{ground}(t1) \wedge \neg w)$, w fresh respectively. They are not to be confused with the WEAKEN function from the paper.

`State::negative_any` (~line 400) implements the WEAKEN algorithm from the paper. `typeinf_options` (~line 624) implements the MIGRATE algorithm in full.

Bibliography

- J. P. Campora, S. Chen, M. Erwig, and E. Walkingshaw. Migrating gradual types. *Proceedings of the ACM on Programming Languages (PACMPL)*, 2(POPL), Dec. 2018.
- A. Kuhlenschmidt, D. Almahallawi, and J. G. Siek. Toward efficient gradual typing for structural types via coercions. In *ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*, 2019.
- Z. Migeed and J. Palsberg. What is decidable about gradual types? *Proceedings of the ACM on Programming Languages (PACMPL)*, 4(POPL), Dec. 2020.
- L. Phipps-Costin, C. J. Anderson, M. Greenberg, and A. Guha. Solver-based gradual type migration. *Proceedings of the ACM on Programming Languages (PACMPL)*, 5(OOPSLA), oct 2021.
- A. Rastogi, A. Chaudhuri, and B. Hosmer. The ins and outs of gradual type inference. In *ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL)*, 2012.
- J. G. Siek and M. Vachharajani. Gradual typing with unification-based inference. In *Dynamic Languages Symposium (DLS)*, 2008.